

# Project Report: Human Activity Recognition using LSTM

## 1. Problem Statement

The goal of this project is to develop a model that can accurately classify human physical activities based on time-series data collected from smartphone sensors. Given sequences of sensor readings from an accelerometer and a gyroscope, the task is to classify the activity being performed into one of six categories: Walking, Walking Upstairs, Walking Downstairs, Sitting, Standing, or Laying. This is a multi-class time-series classification problem that has applications in healthcare monitoring, fitness tracking, and context-aware computing.

## 2. Dataset Description

- **Source:** The project utilizes the **UCI Human Activity Recognition (HAR) Using Smartphones Dataset**.
- **Size and Features:** The dataset is pre-split into a training set and a test set.
  - **Training Set:** 7,352 samples.
  - **Test Set:** 2,947 samples.
  - Each sample consists of a sequence of **128 timesteps**, with **9 features** per timestep. The features are derived from the raw tri-axial signals from the accelerometer and gyroscope and include total acceleration, body acceleration, and body gyroscope readings for the X, Y, and Z axes. The input data has a shape of (samples, 128, 9).
- **Labels:** There are 6 distinct activity classes to be predicted.
- **Preprocessing:** The integer labels (1-6) were first converted to a zero-based index (0-5). Subsequently, they were **one-hot encoded** into a binary vector format (e.g., SITTING as [0,0,0,1,0,0]) to be compatible with the model's output layer and the categorical cross-entropy loss function.

## 3. Methodology

A deep learning approach was chosen to automatically learn temporal patterns from the raw time-series data without extensive manual feature engineering.

- **Algorithm:** A **Long Short-Term Memory (LSTM)** neural network was selected. LSTMs are a type of Recurrent Neural Network (RNN) specifically designed to handle long-term dependencies in sequence data, making them ideal for analyzing sensor readings over time.
- **Rationale:** The model employs a **stacked LSTM architecture**, where multiple LSTM layers are placed one after another. This allows the network to learn hierarchical representations of the time-series data, with earlier layers capturing low-level temporal features and later layers combining them into more abstract patterns related to specific activities.

## 4. Architecture Description

The model was constructed using the TensorFlow Keras Sequential API. It consists of two stacked LSTM layers followed by a final Dense output layer.

The architecture is detailed in the summary table below:

Layer (type)	Output Shape	Param #
lstm_6 (LSTM)	(None, 128, 32)	5,376
lstm_7 (LSTM)	(None, 32)	8,320
dense_3 (Dense)	(None, 6)	198

- **Input Layer:** Implicitly defined with an input shape of (128, 9).
- **First LSTM Layer:** Contains 32 units. The `return_sequences=True` argument ensures that the output of this layer is the full sequence of hidden states, which is necessary for feeding into the subsequent LSTM layer.
- **Second LSTM Layer:** Contains 32 units. It receives the sequence from the first layer and outputs only the hidden state of the final timestep, effectively providing a summary vector of the entire input sequence.
- **Output (Dense) Layer:** A fully connected layer with 6 neurons, corresponding to the 6 activity classes. It uses a **softmax** activation function to output a probability distribution over the classes.

## 5. Evaluation & Implementation

- **Dataset Split:** The pre-defined training set (7,352 samples) was used for training the model. The test set (2,947 samples) was used as the **validation set** during training to monitor performance and as the final hold-out set for evaluation.
- **Loss Function:** **Categorical Crossentropy** was used as the loss function, which is standard for multi-class classification problems with one-hot encoded labels.
- **Evaluation Metrics:** Model performance was primarily assessed using **accuracy**. For a more detailed analysis, **Precision**, **Recall**, and **F1-Score** (with a 'weighted' average) were also calculated on the test set.
- **Implementation Details:**
  - **Optimizer:** Adam
  - **Batch Size:** 64
  - **Epochs:** The model was trained for a maximum of 100 epochs.
- **Other Methods:**
  - **Callbacks:** Two Keras callbacks were used to improve training efficiency and prevent overfitting:
    1. **EarlyStopping:** Monitored the validation loss (`val_loss`) and stopped training if there was no improvement for 5 consecutive epochs. It was configured to restore the weights from the epoch with the best `val_loss`.
    2. **ReduceLROnPlateau:** Reduced the learning rate by a factor of 0.2 if the

validation loss did not improve for 5 epochs.

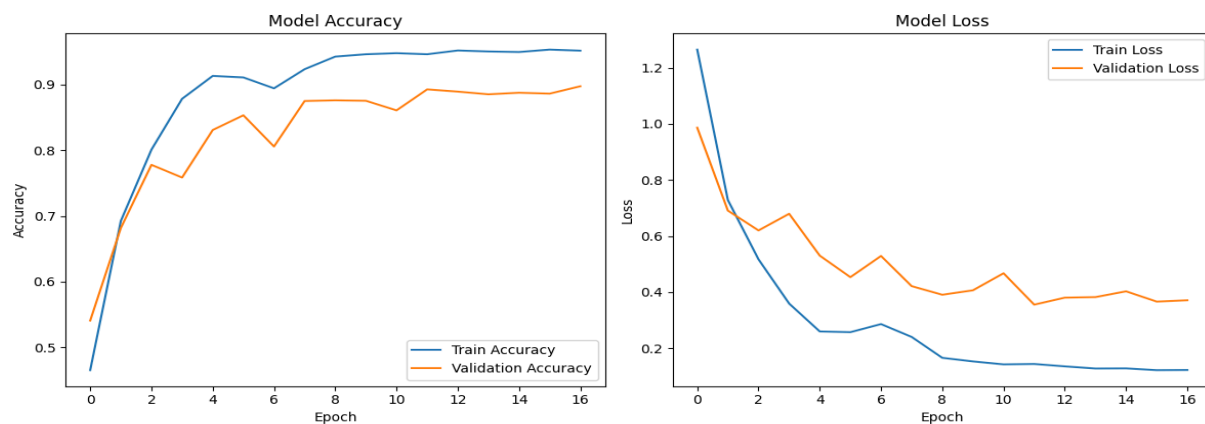
## 6. Results & Analysis

The model achieved a **final test accuracy of 89.24%** with a **test loss of 0.3553**.

### Detailed Classification Metrics:

- **Precision:** 0.8934
- **Recall:** 0.8924
- **F1-Score:** 0.8922

### Visualizations & Analysis:



## 7. Future Improvements

While the model achieved a strong baseline performance, several avenues for improvement exist:

- **Hyperparameter Tuning:** Systematically tune hyperparameters like the number of LSTM units, batch size, and the Adam optimizer's learning rate might help in improving further performance.